



NYU

**DATA
SCIENCE**

word2vec and friends

<https://github.com/bmtgoncalves/word2vec-and-friends/>

Bruno Gonçalves

www.bgoncalves.com

Teaching machines to read!

- Computers are really good at crunching numbers but not so much when it comes to words.
- Perhaps can we represent words numerically?

a	1
about	2
above	3
after	4
again	5
against	6
all	7
am	8
an	9
and	10
any	11
are	12
aren't	13
as	14
...	...

Teaching machines to read!

- Computers are really good at crunching numbers but not so much when it comes to words.
- Perhaps can we represent words numerically?

$$\begin{aligned}v_{after} &= (0, 0, 0, 1, 0, 0, \dots)^T \\v_{above} &= (0, 0, 1, 0, 0, 0, \dots)^T\end{aligned}\quad \begin{array}{l}\text{One-hot} \\ \text{encoding}\end{array}$$

- Can we do it in a way that preserves **semantic** information?

"You shall know a word by the company it keeps"
(J. R. Firth)



- **Words** that have similar **meanings** are used in similar **contexts** and the context in which a word is used helps us understand its meaning.

The red **house** is beautiful.
The blue **house** is old.
The red **car** is beautiful.
The blue **car** is old.

Teaching machines to read!

"You shall know a word by the company it keeps"
(J. R. Firth)



→ Words with similar meanings should have similar representations.

→ From a word we can get some idea about the context where it might appear

_____ house _____
_____ car _____

$$\max p(C|w)$$

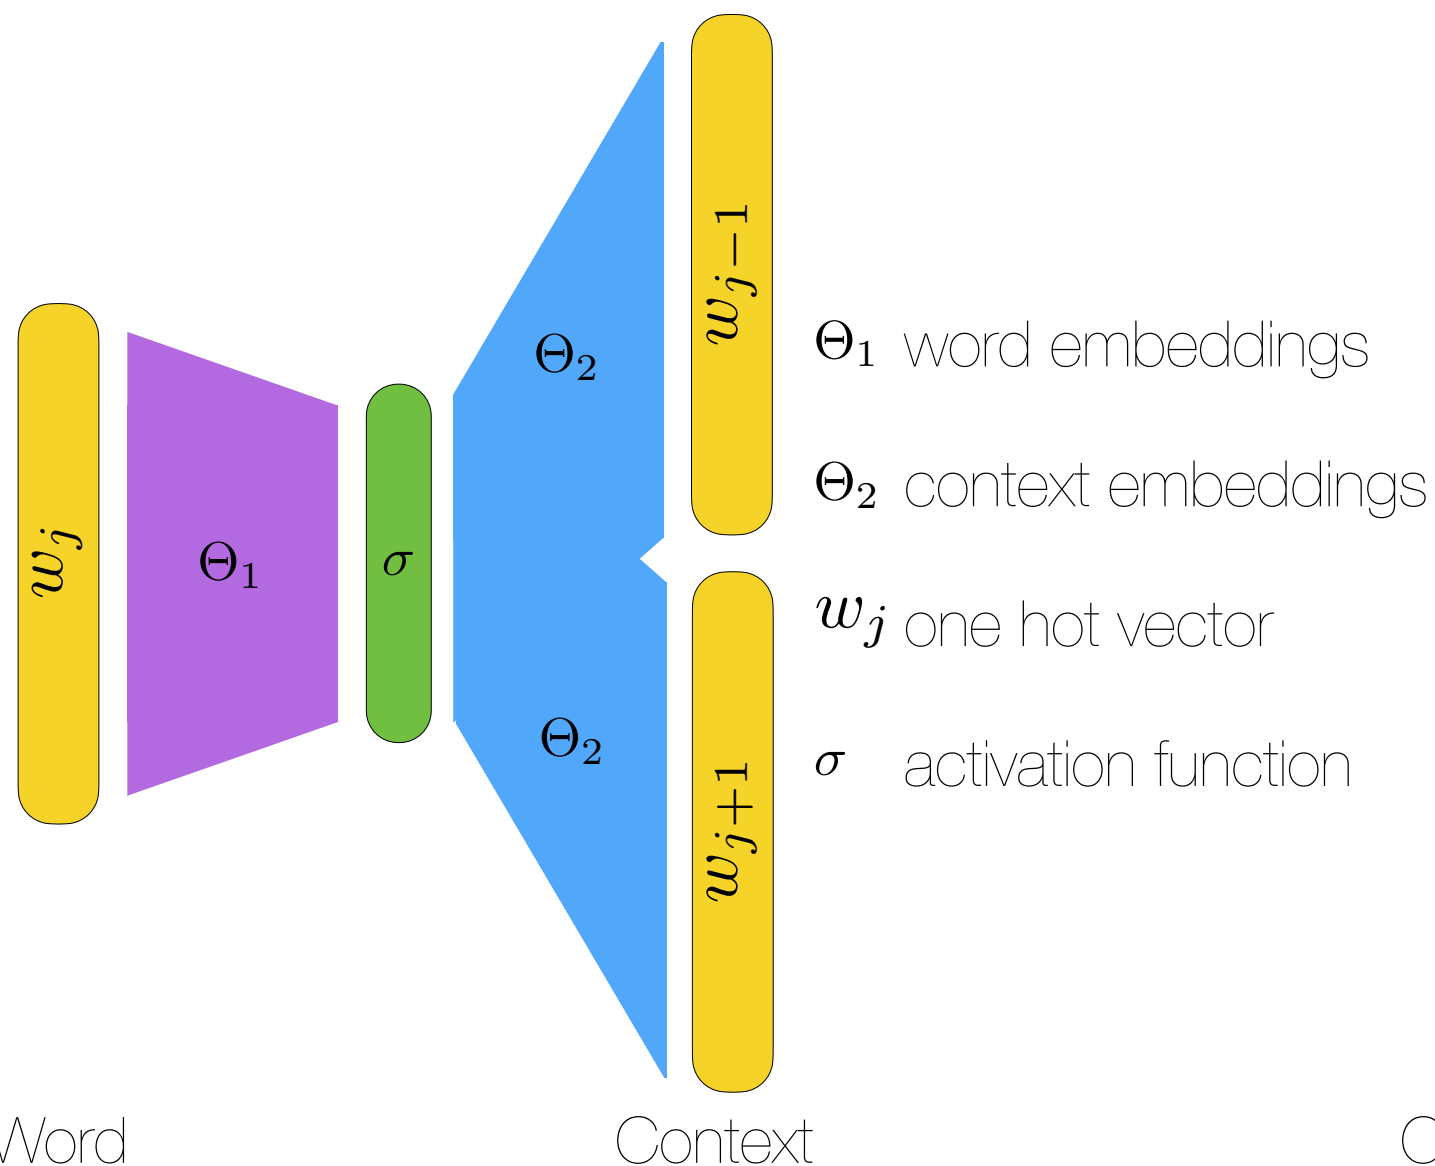
→ And from the context we have some idea about possible words

The red _____ is beautiful.
The blue _____ is old.

$$\max p(w|C)$$


Skipgram

$$\max p(C|w)$$



Continuous Bag of Words

$$\max p(w|C)$$



Skipgram

- Let us take a better look at a simplified case with a single context word
- Words are one-hot encoded vectors $w_j = (0, 0, 1, 0, 0, 0, \dots)^T$ of length V


- Θ_1 is an $(M \times V)$ matrix so that when we take the product:

$$\Theta_1 \cdot w_j$$

- We are effectively selecting the j 'th column of Θ_1 :

$$v_j = \Theta_1 \cdot w_j$$

- The **linear** activation function simply passes this value along which is then multiplied by Θ_2 , a $(V \times M)$ matrix.



- Each element k of the output layer its then given by:

$$u_k^T \cdot v_j$$

- We convert these values to a normalized probability distribution by using the **softmax**

Softmax

- A standard way of converting a set of number to a normalized probability distribution:

$$\text{softmax}(x) = \frac{\exp(x_j)}{\sum_l \exp(x_l)}$$

- With this final ingredient we obtain:

$$p(w_k | w_j) \equiv \text{softmax}(u_k^T \cdot v_j) = \frac{\exp(u_k^T \cdot v_j)}{\sum_l \exp(u_l^T \cdot v_j)}$$


- Our goal is then to learn:

$$\Theta_1 \quad \Theta_2$$

- so that we can predict what the next word is likely to be using:

$$p(w_{j+1} | w_j)$$

- But how can we quantify how far we are from the correct answer? Our error measure shouldn't be just binary (right or wrong)...



Cross-Entropy

- First we have to recall that what we are, in effect, comparing two probability distributions:

$$p(w_k|w_j)$$

- and the one-hot encoding of the context:

$$w_{j+1} = (0, 0, 0, 1, 0, 0, \dots)^T$$

- The Cross Entropy measures the distance, in number of bits, between two probability distributions \mathbf{p} and \mathbf{q} :

$$H(p, q) = - \sum_k p_k \log q_k$$

- In our case, this becomes:

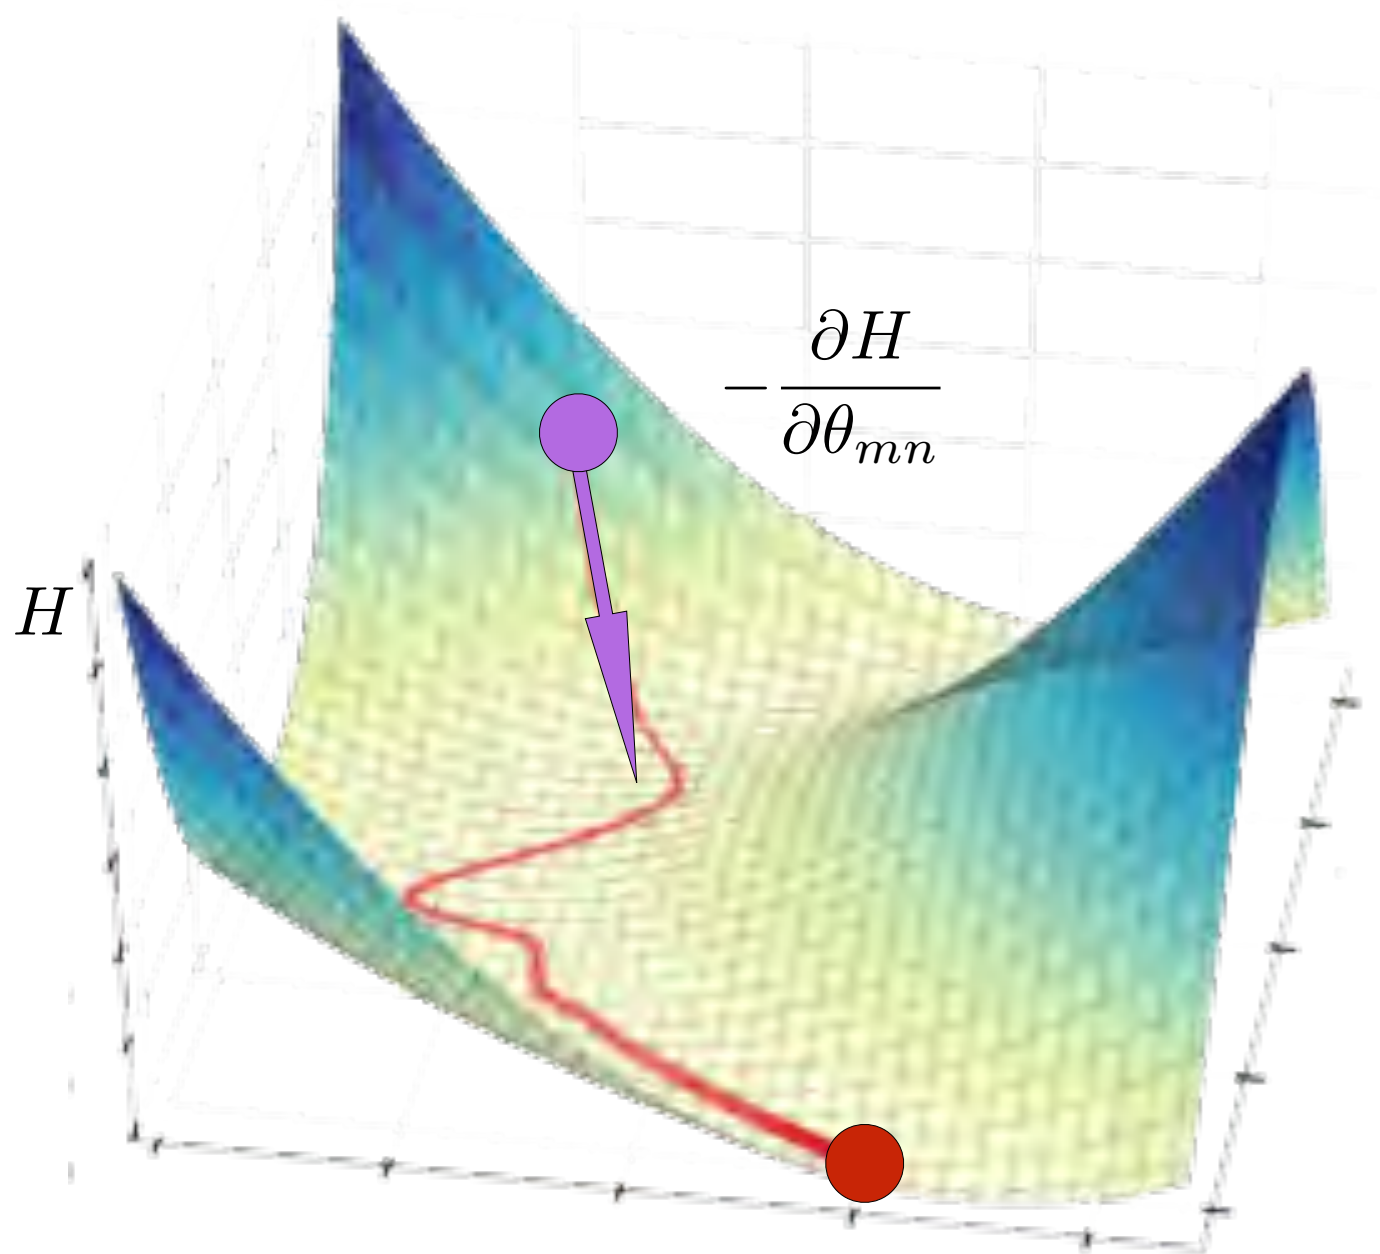
$$H[w_{j+1}, p(w_k|w_j)] = - \sum_k w_{j+1}^k \log p(w_k|w_j)$$

- So it's clear that the only non zero term is the one that corresponds to the “hot” element of w_{j+1}

$$H = - \log p(w_{j+1}|w_j)$$

- This is our Error function. But how can we use this to update the values of Θ_1 and Θ_2 ?

Gradient Descent



- Find the gradient for each training batch
- Take a step **downhill** along the direction of the gradient

$$\theta_{mn} \leftarrow \theta_{mn} - \alpha \frac{\partial H}{\partial \theta_{mn}}$$

- where α is the step size.
- Repeat until “convergence”.

Chain-rule

- How can we calculate

$$\frac{\partial H}{\partial \theta_{mn}} = \frac{\partial}{\partial \theta_{mn}} \log p(w_{j+1}|w_j) \quad \theta_{mn} = \left\{ \theta_{mn}^{(1)}, \theta_{mn}^{(2)} \right\}$$

- we rewrite:

$$\frac{\partial H}{\partial \theta_{mn}} = \frac{\partial}{\partial \theta_{mn}} \log \frac{\exp(u_k^T \cdot v_j)}{\sum_l \exp(u_l^T \cdot v_j)}$$

- and expand:

$$\frac{\partial H}{\partial \theta_{mn}} = \frac{\partial}{\partial \theta_{mn}} \left[u_k^T \cdot v_j - \log \sum_l \exp(u_l^T \cdot v_j) \right]$$

- Then we can rewrite:

$$u_k^T \cdot v_j = \sum_q \theta_{kq}^{(2)} \theta_{qj}^{(1)}$$

- and apply the chain rule:


$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \frac{\partial g(x)}{\partial x}$$

Training procedures



- **online learning** - update weights after **each** case
 - might be useful to update model as new data is obtained
 - subject to fluctuations
- **mini-batch** - update weights after a "**small**" number of cases
 - batches should be balanced
 - if dataset is redundant, the gradient estimated using only a fraction of the data is a good approximation to the full gradient.
- **momentum** - let gradient change the **velocity** of weight change instead of the value directly
- **rmsprop** - divide learning rate for each weight by a **running average** of "recent" gradients
- **learning rate** - vary over the course of the training procedure and use different learning rates for each weight

SkipGram with Larger Contexts

$$H = -\log p(w_{j+1}|w_j)$$




$$H = -\frac{1}{T} \sum_t \log p(w_{j+t}|w_j)$$



- Use the same Θ_2 for all context words.
- Use the average of cross entropy.
- word order is not important (the average does not change)
- Can essentially be trained **one context word at a time**..

Continuous Bag of Words

- The process is essentially the same



Variations

- Hierarchical Softmax:


- Approximate the **softmax** using a binary tree
- Reduce the number of calculations per training example from V to $\log_2 V$ and increase performance by orders of magnitude.

- Negative Sampling:

- Under sample the most frequent words by removing them from the text **before** generating the contexts
- Similar idea to removing **stop-words** — very frequent words are less informative.
- Effectively makes the window larger, increasing the amount of information available for context

Comments


- **word2vec**, even in its original formulation is actually a family of algorithms using various combinations of:
 - Skip-gram, CBOW
 - Hierarchical Softmax, Negative Sampling
- The output of this neural network is deterministic:
 - If two words appear in the same context ("blue" vs "red", for e.g.), they will have similar internal representations in Θ_1 and Θ_2
 - Θ_1 and Θ_2 are vector embeddings of the input words and the context words respectively
- Words that are too rare are also removed.
- The original implementation had a dynamic window size:
 - for each word in the corpus a window size k' is sampled uniformly between 1 and k




Online resources

- C - <https://code.google.com/archive/p/word2vec/> (the original one)
- Python/tensorflow - <https://www.tensorflow.org/tutorials/word2vec>
 - Both a minimalist and an efficient versions are available in the tutorial
- Python/gensim - <https://radimrehurek.com/gensim/models/word2vec.html>
- Pretrained embeddings:
 - 30+ languages, <https://github.com/Kyubyong/wordvectors>
 - 100+ languages trained using wikipedia: <https://sites.google.com/site/rmyeid/projects/polyglot>


Visualization



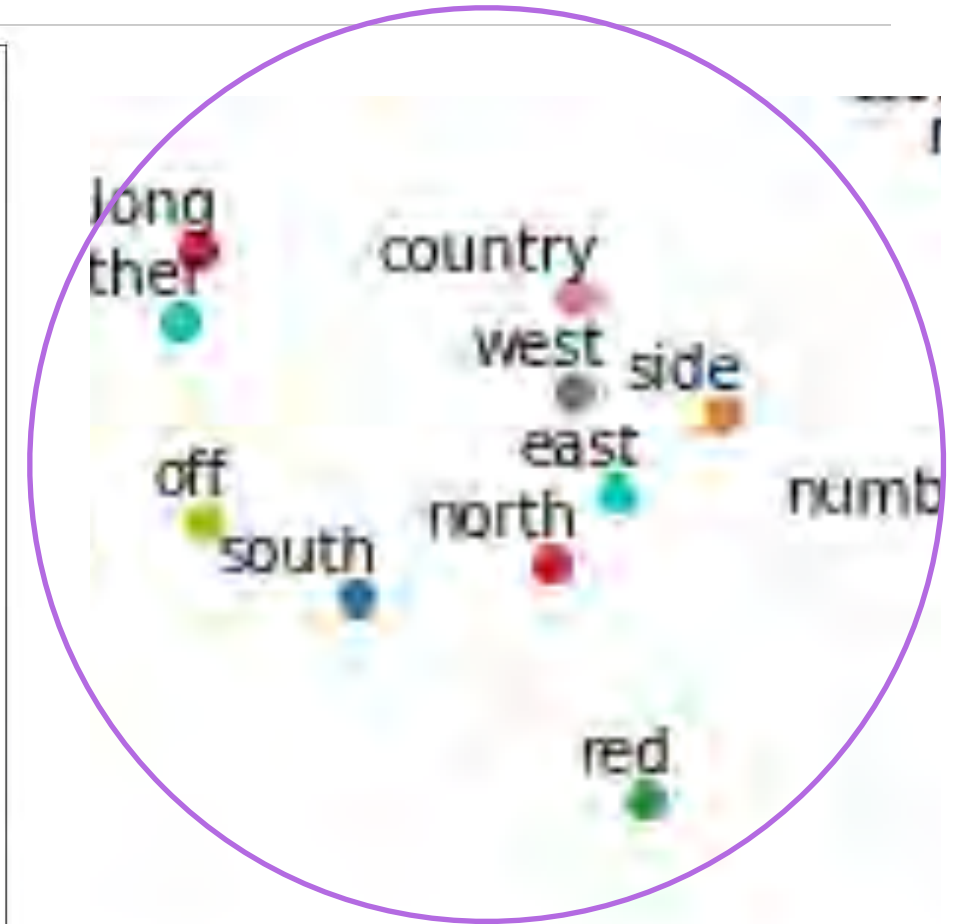

Visualization



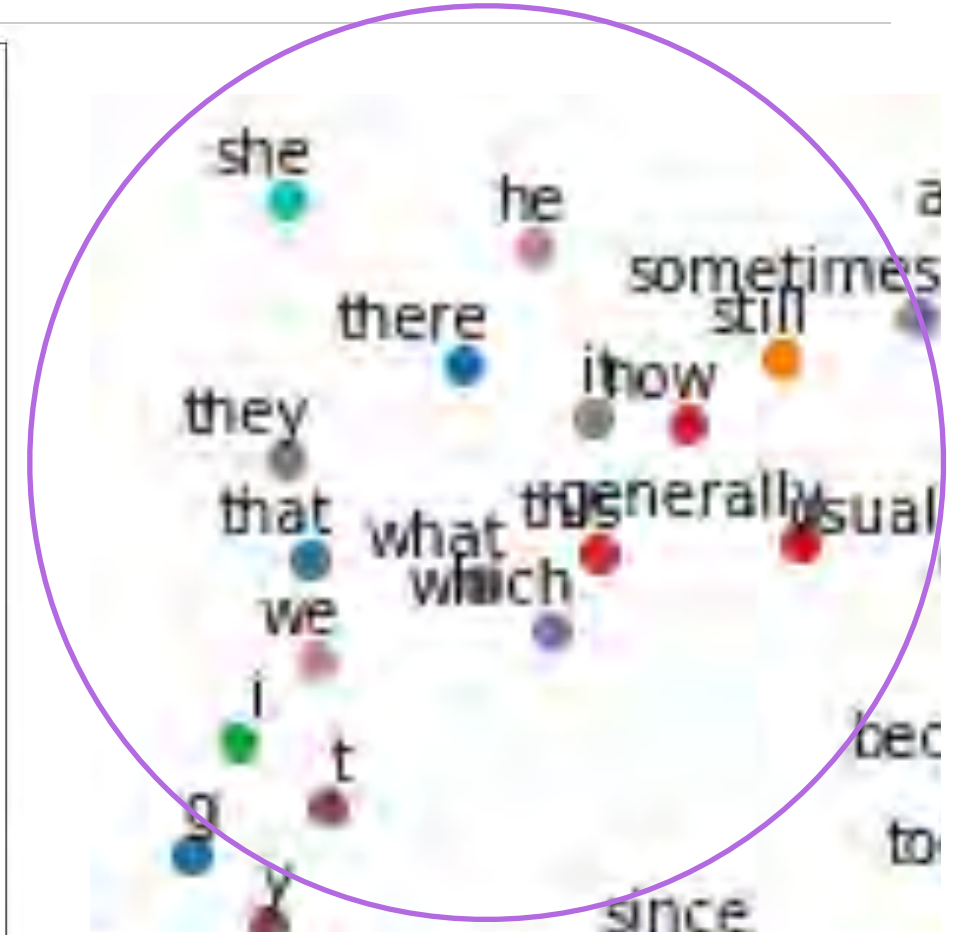

Visualization




Visualization



Visualization



Visualization



Analogies

"You shall know a word by the company it keeps"
(J. R. Firth)



- The embedding of each word is a function of the context it appears in:

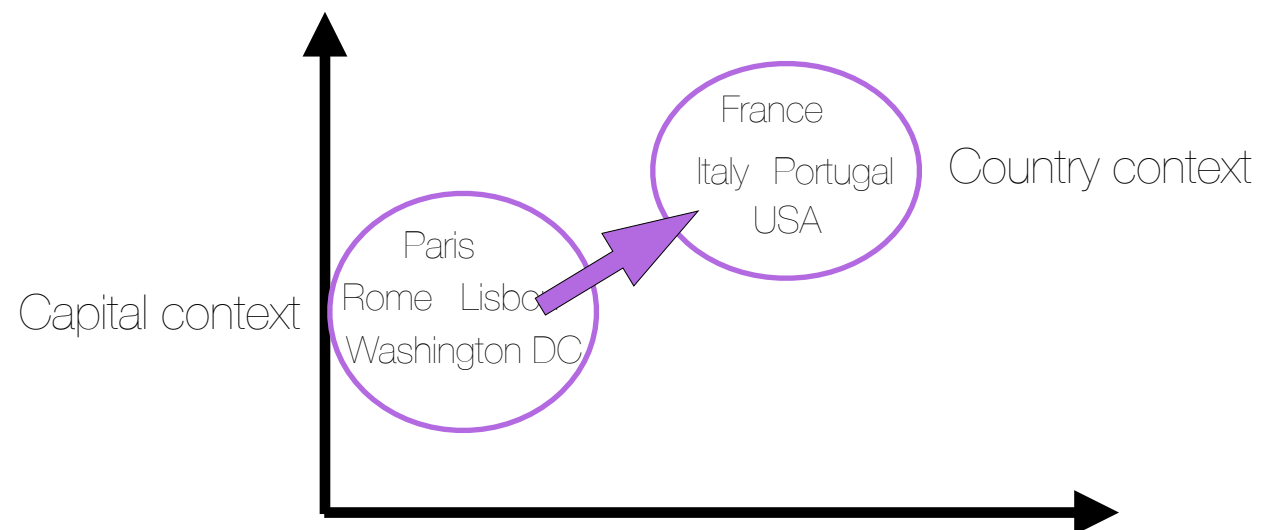
$$\sigma(\text{red}) = f(\text{context}(\text{red}))$$

- words that appear in similar contexts will have similar embeddings:

$$\text{context}(\text{red}) \approx \text{context}(\text{blue}) \implies \sigma(\text{red}) \approx \sigma(\text{blue})$$

- "Distributional hypothesis" in linguistics

Geometrical relations
between contexts imply
semantic relations
between words!



$$\sigma(\text{France}) - \sigma(\text{Paris}) + \sigma(\text{Rome}) = \sigma(\text{Italy})$$

$$\vec{b} - \vec{a} + \vec{c} = \vec{d}$$

Analogies


<https://www.tensorflow.org/tutorials/word2vec>

$$\vec{b} - \vec{a} + \vec{c} = \vec{d}$$


$$d^\dagger = \operatorname{argmax}_x \frac{(\vec{b} - \vec{a} + \vec{c})^T}{\|\vec{b} - \vec{a} + \vec{c}\|} \vec{x}$$

$$d^\dagger \sim \operatorname{argmax}_x (\vec{b}^T \vec{x} - \vec{a}^T \vec{x} + \vec{c}^T \vec{x})$$


What is the word **d** that is most **similar** to **b** and **c** and most **dissimilar** to **a**?



Male-Female



Verb tense

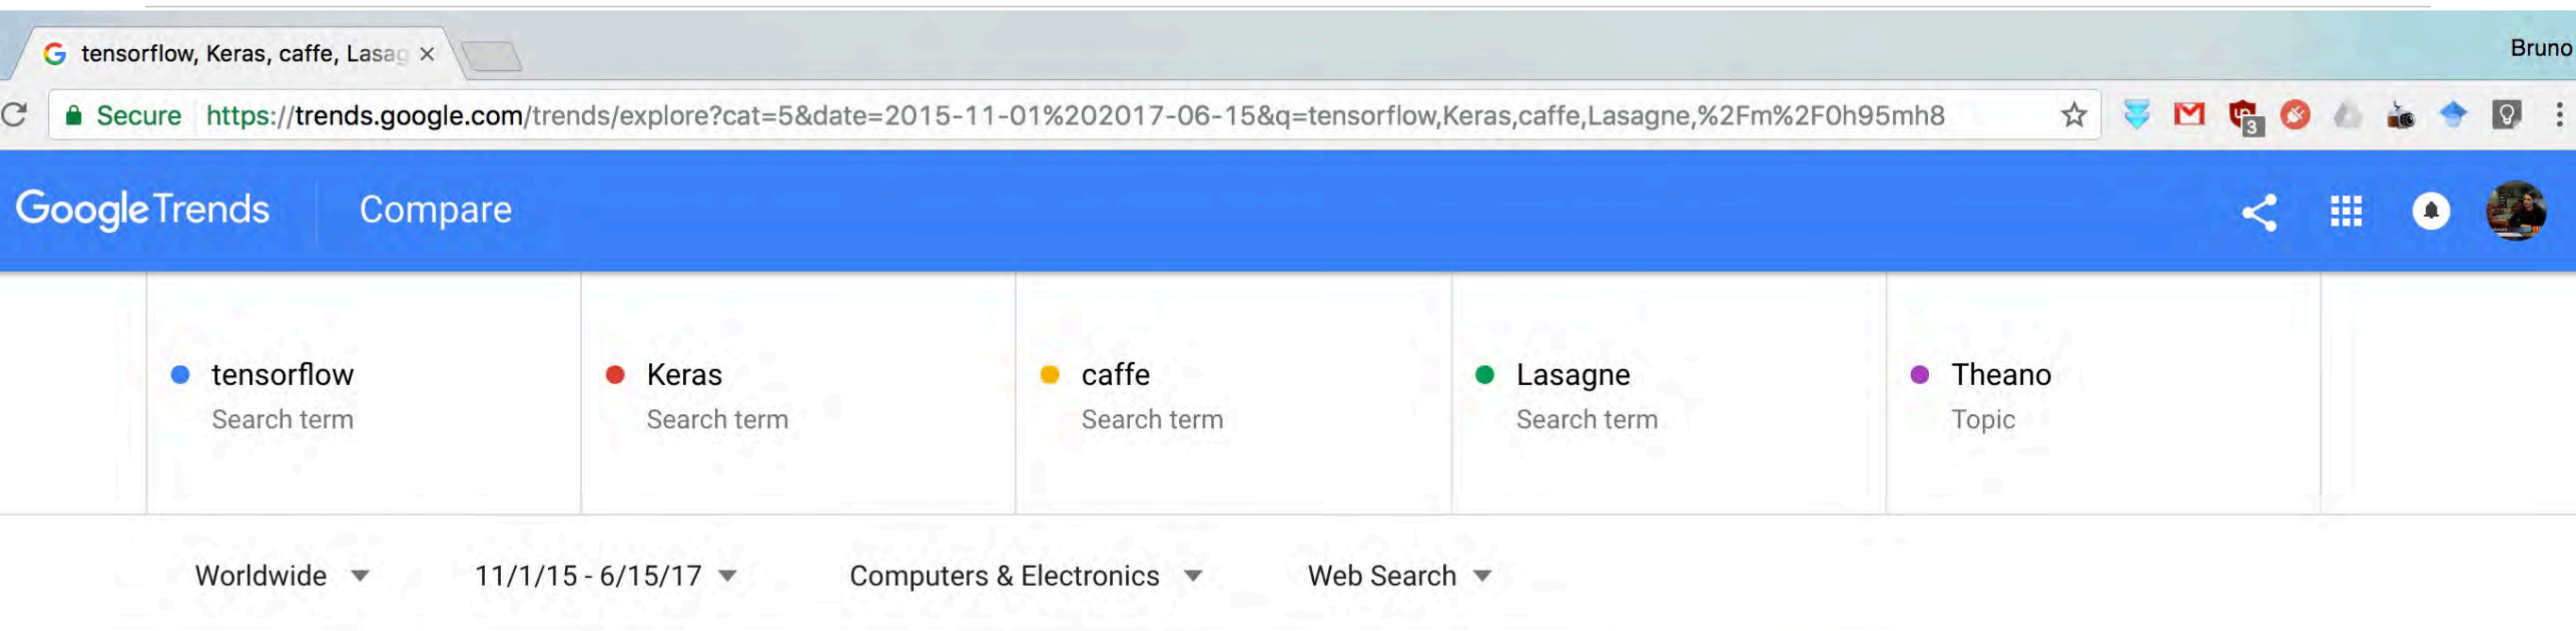


Country-Capital




<https://github.com/bmtgoncalves/word2vec-and-friends/>

Tensorflow



Search terms match specific words; topics are concepts that match similar terms in any language. [Learn more](#)




A diversion...

- Let's imagine I want to perform these calculations:

$$y = f(x)$$

$$z = g(y)$$

- for some given x .
- To calculate z we must follow a certain sequence of operations.




A diversion...

- Let's imagine I want to perform these calculations:

$$y = f(x)$$


$$z = g(y)$$

- for some given x .
- To calculate z we must follow a certain sequence of operations.
- Which can be shortened if we are interested in just the value of y
- In **Tensorflow**, this is called a **Computational Graph** and it's the most fundamental concept to understand
- Data, in the form of **tensors**, **flows** through the graph from inputs to outputs
- Tensorflow**, is, essentially, a way of defining arbitrary computational graphs in a way that can be automatically distributed and optimized.




Computational Graphs

- If we use base functions, tensorflow knows how to automatically calculate the respective gradients
- Automatic BackProp
- Graphs can have multiple outputs
 - Predictions
 - Cost functions
 - etc...



Sessions

- After we have defined the computational graph, we can start using it to make calculations
- All computations must take place within a “session” that defines the values of all **required** input values
- Which values are required for a specific computation depend on what part of the graph is actually being executed.
- When you request the value of a specific output, tensorflow determines what is the specific **subgraph** that must be executed and what are the required input values.
- For optimization purposes, it can also execute independent parts of the graph in different devices (CPUs, GPUs, TPUs, etc) at the same time.



Installation

<https://www.tensorflow.org/install/>

Install TensorFlow

Assuming the prerequisite software is installed on your Mac, take the following steps:

1. Install TensorFlow by invoking **one** of the following commands:

```
$ pip install tensorflow      # Python 2.7; CPU support
$ pip3 install tensorflow     # Python 3.n; CPU support
```

If the preceding command runs to completion, you should now [validate your installation](#).

A basic Tensorflow program

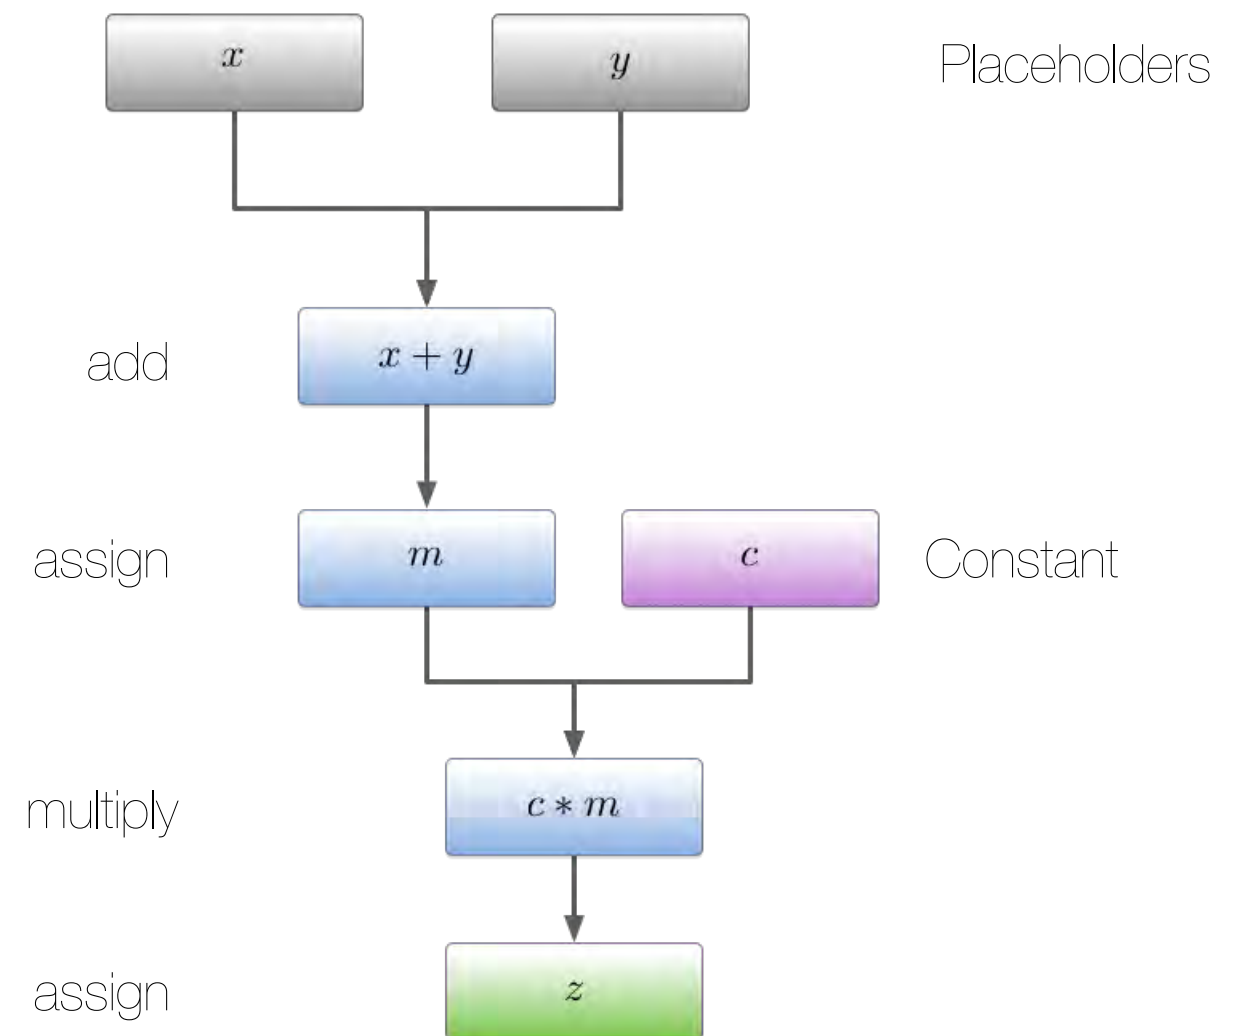
$$z = c * (x + y)$$

```
import tensorflow as tf
```

```
x = tf.placeholder(tf.float32)  
y = tf.placeholder(tf.float32)  
c = tf.constant(3.)
```

```
m = tf.add(x, y)  
z = tf.multiply(m, c)
```

```
with tf.Session() as sess:  
    output = sess.run(z, feed_dict={x: 1., y: 2.})  
    print("Output value is:", output)
```



A basic Tensorflow program

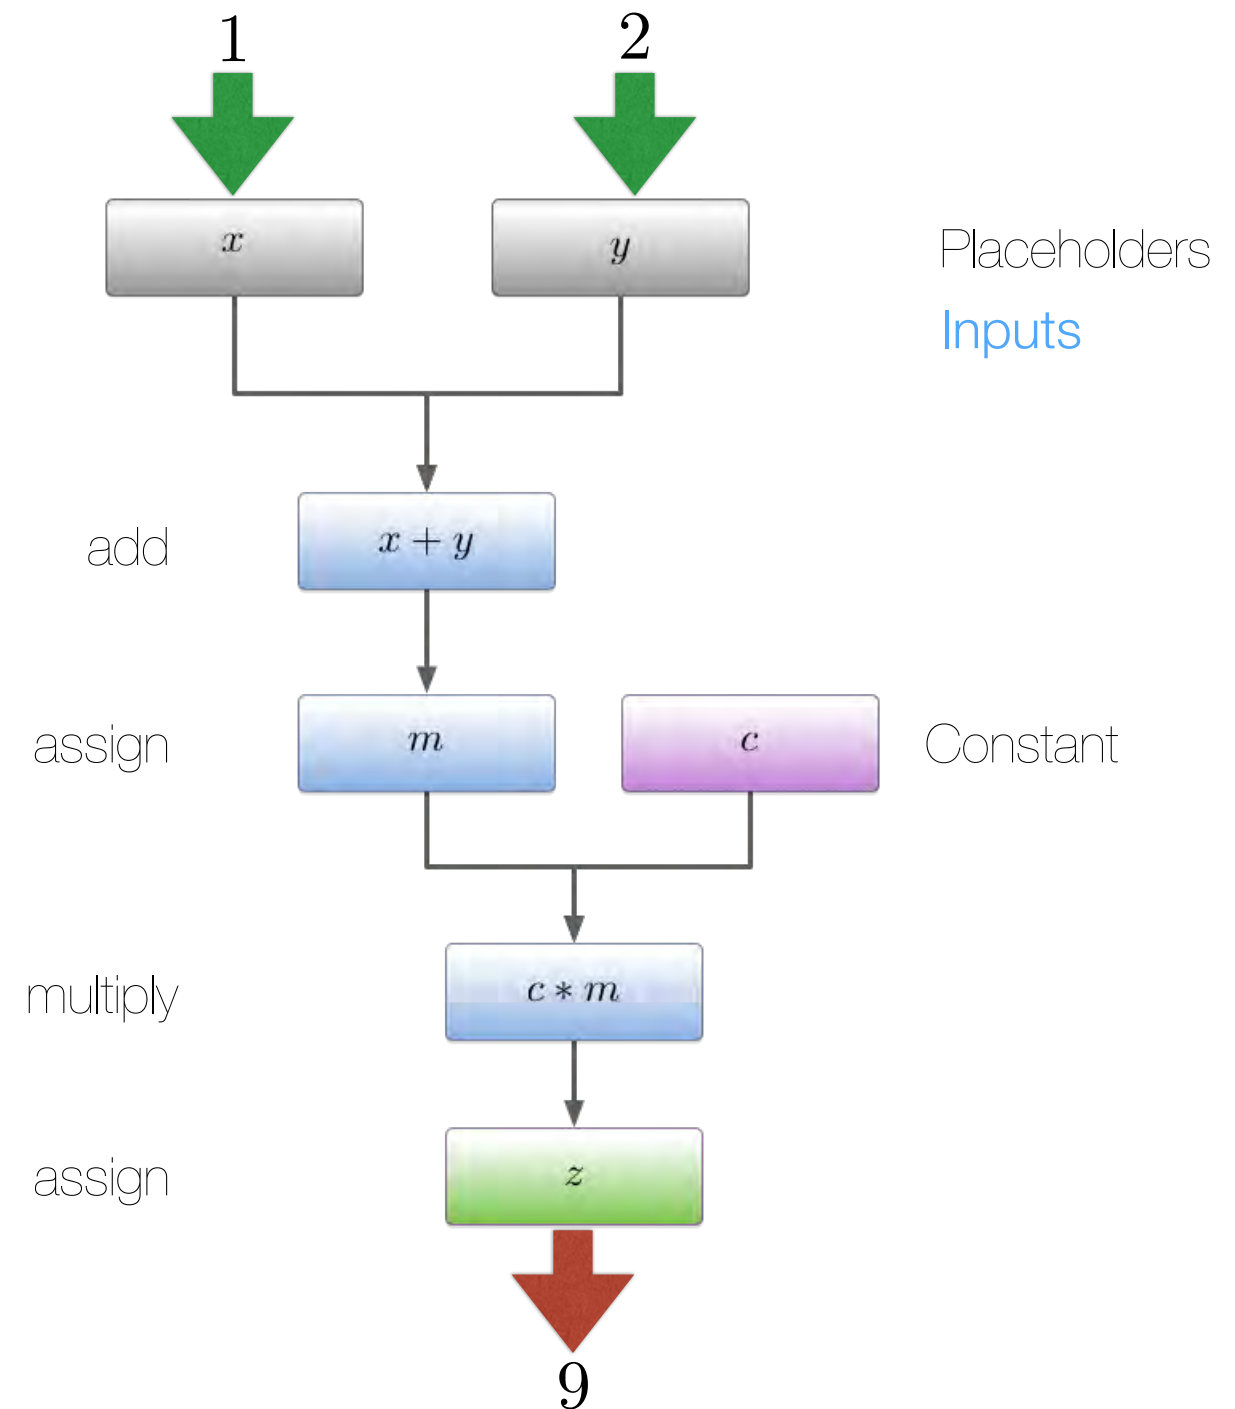
$$z = c * (x + y)$$

```
import tensorflow as tf

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
c = tf.constant(3.)

m = tf.add(x, y)
z = tf.multiply(m, c)

with tf.Session() as sess:
    output = sess.run(z, feed_dict={x: 1., y: 2.})
    print("Output value is:", output)
```



Jupyter Notebook

Statistically Significant Detection of Linguistic Change


Vivek Kulkarni
Stony Brook University, USA
vvkulkarni@cs.stonybrook.edu

Bryan Perozzi
Stony Brook University, USA
bperozzi@cs.stonybrook.edu

Rami Al-Rfou
Stony Brook University, USA
ralrfou@cs.stonybrook.edu

Steven Skiena
Stony Brook University, USA
skiena@cs.stonybrook.edu

- Train word embeddings for different years using Google Books
- Independently trained embeddings differ by an arbitrary rotation
- Align the different embeddings for different years
- Track the way in which the meaning of words shifted over time!




node2vec: Scalable Feature Learning for Networks

Aditya Grover
Stanford University
adityag@cs.stanford.edu


Jure Leskovec
Stanford University
jure@cs.stanford.edu

- You can generate a graph out of a sequence of words by assigning a node to each word and connecting the words within their neighbors through edges.
- With this representation, a piece of text is a walk through the network. Then perhaps we can invert the process? Use walks through a network to generate a sequence of nodes that can be used to train node embeddings?
- node embeddings should capture **features** of the network structure and allow for detection of similarities between nodes.

- The features depends strongly on the way in which the network is traversed
- Generate the contexts for each node using Breath First Search and Depth First Search



- Perform a biased Random Walk



- BFS - Explores only limited neighborhoods. Suitable for structural equivalences
- DFS - Freely explores neighborhoods and covers homophiles communities
- By modifying the parameter of the model it can interpolate between the BFS and DFS extremes

dna2vec: Consistent vector representations of
variable-length k-mers

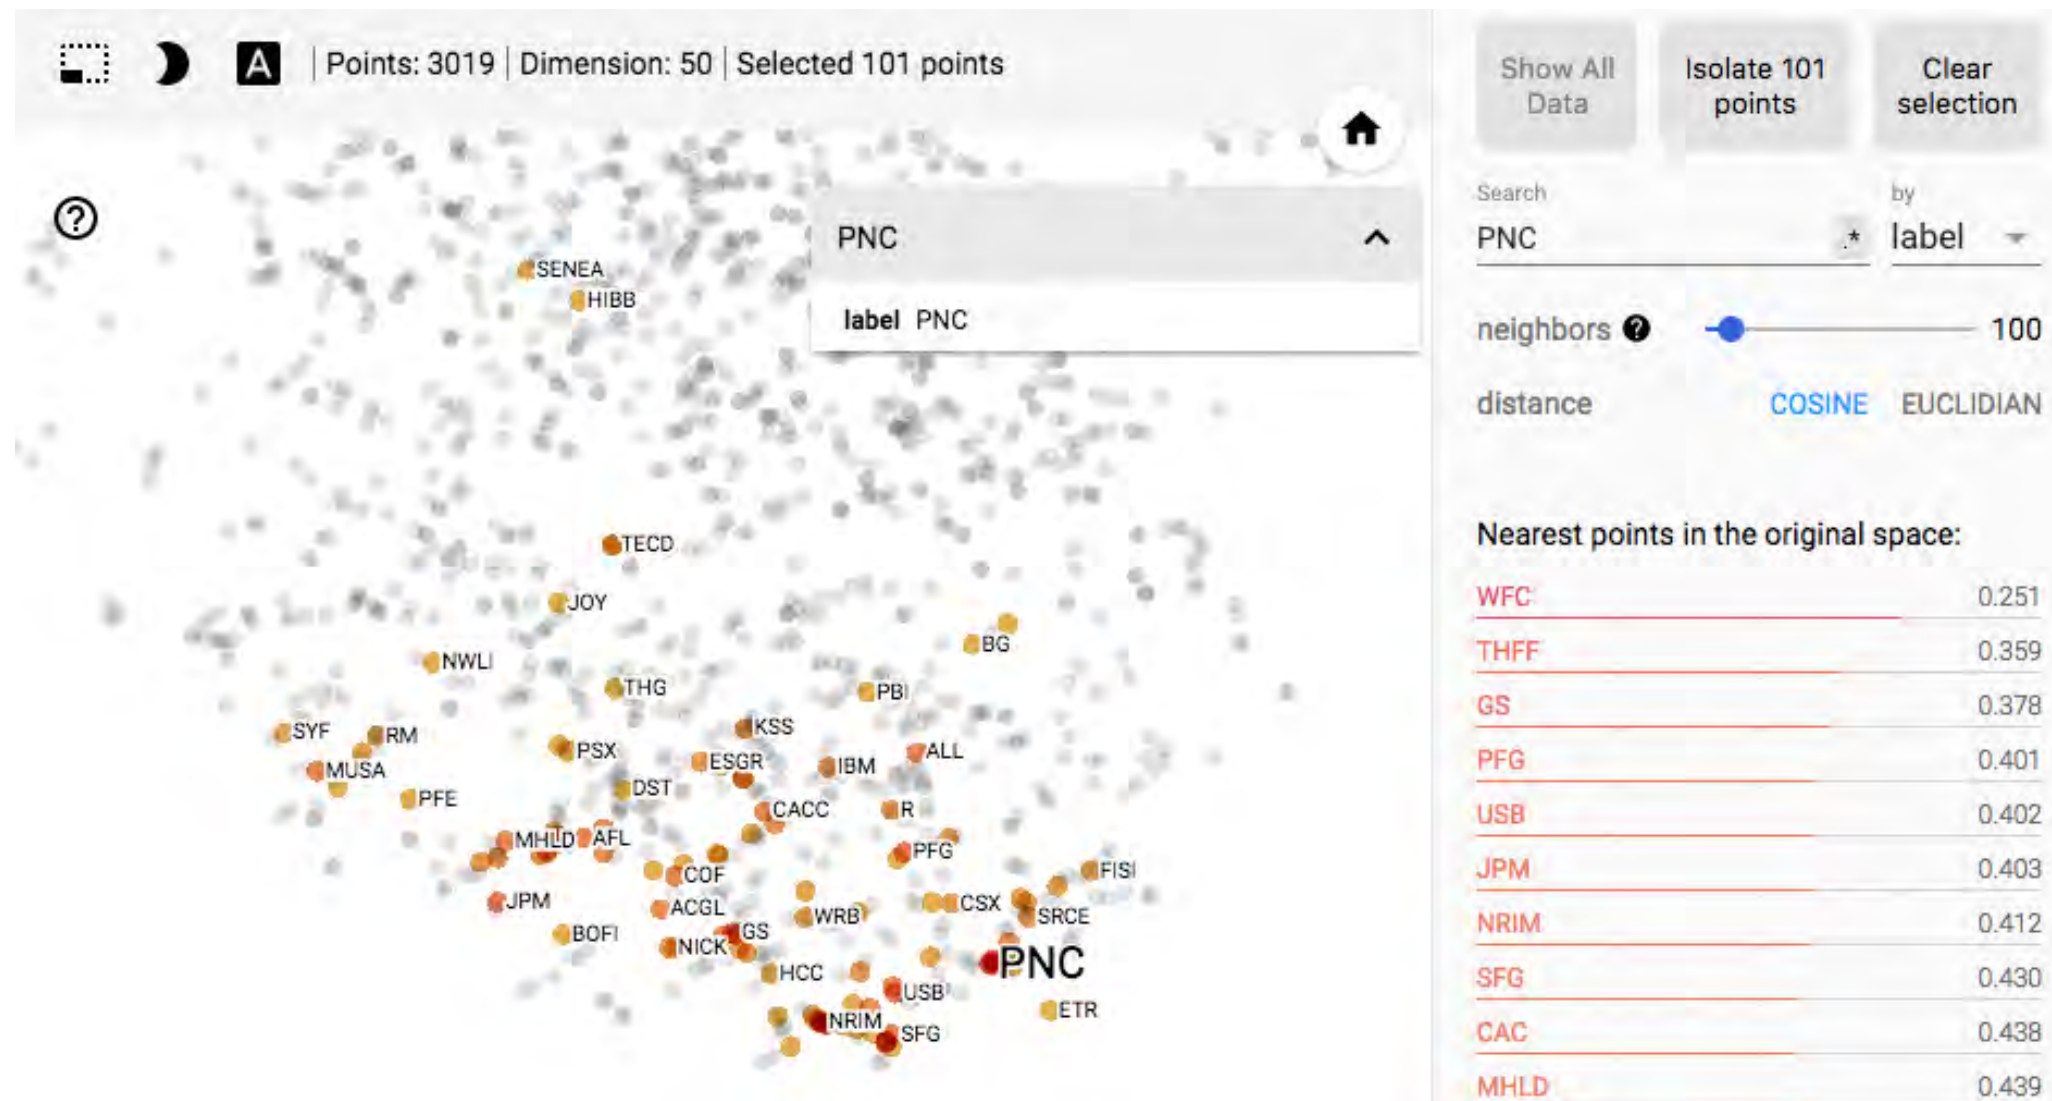
Patrick Ng
ppn3@cs.cornell.edu

- Separate the genome into long non-overlapping DNS fragments.
- Convert long DNA fragments into overlapping **variable** length k-mers
- Train embeddings of each k-mer using **Gensim** implementation of SkipGram.
 - Summing embeddings is related to concatenating **k-mers**
 - Cosign similarity of k-mer embeddings reproduces a biologically motivated similarity score (Needleman-Wunsch) that is used to align nucleoti

stock2vec


<https://medium.com/towards-data-science/stock2vec-from-ml-to-p-e-2e6ba407c24>

- Apply word2vec to the 40 years of stock market data
- Identify significant semantic similarities between companies working in the same area



stock2

ml-to-p-e-2e6ba407c24



@bgoncalves

igoncalves.com

Thank you!

You can hear me speak more about
word2vec in this weeks podcast!

